

Deep Active Learning In The Presence Of Label Noise

Moseli mots'oehli^[0000-1111-2222-3333] and
Kyungim Baek^[1111-2222-3333-4444]

Department of Information and Computer Science,
University of Hawai'i at Manoa
{moselim, kyungim}@hawaii.edu

1 Introduction

Machine learning algorithms are a sub-class of artificial intelligence that learn from data to perform a pre-defined task such as classification, regression, or clustering. Of the numerous algorithms for machine learning, artificial neural networks, deep neural networks in particular have done exceptionally well in tasks involving complex data representations such as images, text, and sound. The main reason for this was the discovery that if you have a large enough dataset, you can build more extensive and more complex models with little to no risk of over-fitting. While this works in theory, the practical applications have major drawbacks such as the need for labeled training examples that come at a high cost due to the time needed to label the data, the high cost of labor in very specialized fields, or the cost of running simulations that would produce the ground truth dataset. The solution comes in the form of Deep active learning (DAL) algorithms, which strive to let the learning algorithm iteratively pick data examples to be labeled from a larger unlabelled dataset, in such a manner that results in (1) a smaller labeled training set, (2) that is representative of the underlying data distribution leading to a near-optimal learner, (3) while at the same time not exceeding the labeling budget. While this works well for most use cases, Real industry dataset labelling has inherent label noise due to a variety of factors such as redundant observations being labelled differently, the best human expert classification performance being low or use of auto-labelling software such as Mechanical Turk. This has adverse effects on these DAL algorithm's performance, and most existing DAL literature focus on noise free settings. We explore existing literature around the problem of using DAL algorithms in the existence of label noise. We are particularly interested in the image detection, segmentation and classification domain using different Deep representation learning frameworks such as Deep Neural networks, convolutional neural networks (CNNs), and vision transformer networks. We conclude by exploring possible directions for future research in DAL in vision tasks under label noise.

2 Preliminary

in this section, we briefly describe Deep Learning (DL), Active learning (AL), and set the scene for label noise in DAL.

2.1 Deep Learning

Deep Learning (DL) refers to the use of Artificial Neural Networks with multiple hidden layers [30] to approximate known or unknown functions $f: X \mapsto Y$. Over the years, different domain specific DL architectures have been developed to enhance the quality of the learned representations from the different data Modalities. Early research focused on improving optimization, custom layers and connections, activation functions, loss functions and hyper-parameter tuning techniques for the multi-layer perceptron as a way to improve performance on different data modalities. For tabular data, tree based ensemble learning algorithms such as Random forest [5], XGBoost [6], and CatBoost [28] are preferred over DL for their superior performance and resource efficiency. A non-exhaustive selection of interesting neural network adaptations to tabular data includes [31,29,27,2,3]. In the Natural language processing domain, earlier work involved learning word and sentence representation using shallow neural networks in an unsupervised setting [24,23,16]. Until the wide adoption of attention based transformer language models [36,32], word and sentence level embeddings would be fed to a deep neural network with recurrent connections such as a Long-Short-Term-Memory(LSTM) network [18] to achieve state-of-the-art results on down-stream text classification, sentence completion, Named entity recognition or summarization tasks. For static visual tasks such as image classification, regression and segmentation, CNN based architectures with specialized layers and prepossessing transformations were eventually superseded by vision transformer models [19].

Since this work focuses on DL algorithms for vision tasks, we explore models for supervised classification, regression, and segmentation. We first give a brief overview of CNNs that are responsible for a large share of progress in vision based tasks. We then highlight the use of hybrid LSTM and CNN networks for tasks with visual and temporal properties, such as is the case for video based classification. Finally we highlight literature on state of the art (S.O.T.A) spatial attention based models in the context of visual tasks (Vision Transformers) [19].

Convolutional Neural Networks

Vision Transformers Before full transformer models in the language domain, the best LSTM models used a low dimensional vector representation to passed information from an encoder network to a decoder network, while using an attention mechanism. Attention in this setting is used to learn what parts of an input sequence are most important when predicting different parts of the output. This method proves to work well, and so was adapted to work with CNNs for image data [11,35].

2.2 Active learning

In most supervised machine learning use cases, there is an initial data collection and labeling cost, in both money and time. In some domains and tasks, datasets are inherently difficult to label for a variety of reasons, meaning more time is needed even by an expert human annotator to assign a label to each sample. In other cases the cost of hiring expert annotators is high, or the cost of producing the samples is high, such as is the case in experimental physics where observations come from very expensive telescopes or particle accelerators. This presents a challenge to the real world use of machine learning systems, especially as unlabeled dataset sizes increase. Active learning is a machine learning paradigm that seeks to address this problem by letting learning algorithms iteratively select a subset L^m of size m , from a larger unlabelled dataset U^n of size $n : m \leq n$, to be labelled by an oracle O for training. The active learning mantra can be stated as follows: *Train a machine learning model on a significantly smaller labelled dataset, with little to no drop in test performance, all the while staying within a pre-determined labelling budget B .*

Active Learning algorithms, while overlapping, can broadly be grouped into: pool based methods, density based methods, and data expansion methods. Pool based methods select samples for labeling from an unlabelled pool, based on either the uncertainty of the current trained model on samples U^{n-m} , the diversity of samples in the labelled set L^m used to train the current model, or a combination of both [20,21,10]. Pool-based methods are simple in their formation and implementation, but can be computationally expensive for large datasets of high dimensional data such as images. Since pool-based methods largely rely on metrics evaluated on the entire unlabeled dataset to select new candidates, they are not ideal for applications that require low latency. Density based methods seek to capture key characteristics of the underlying data distribution. This is done by selecting a core-set of samples for labeling that are sufficiently representative of the entire dataset, and lead to good generalization [33,26,25]. More recent literature blends pool and core-set methods to take advantage of each approach's benefits. These methods thus lead to efficient and robust models trained on core-sets containing diverse samples that maximize the margins between class margins [17,13]. Some methods in this approach use the hidden layer representations from training a self-supervised task on the image data, instead of the raw pixels. These include pre-training on image orientation: random $(90, 180, 270, 360)^\circ$ rotation classification, or self-supervised contrastive learning, where the target is an arbitrary patch of adjacent pixels in the image [7,12,37]. Data expansion methods seek to expand the training dataset, by generating reasonably realistic synthetic data samples for each target class at only a computational cost than that of obtaining human labeled samples, while enhancing the learning algorithm's performance on the real test dataset [?]. Since their introduction, generative adversarial networks (GANs) and their variations, [15,14,34], were the go-to method for generating synthetic data. However, the samples tend to be unrealistic, the training unstable, and lacking intrinsic evaluation metrics [4,22].

2.3 Label Noise

Label Noise refers to the scenario in which data labels are corrupted, with or without intention, so that we do not have 100% confidence in their correctness. Label noise is different from feature noise which is normally used to refer to adding Gaussian noise to feature values. Label noise impacts learning algorithms more adversely than feature noise does, and is harder to deal with [8,38,1,9]. Learning noisy labels is especially hard due to the facts that cost functions are generally significantly less complex than feature extraction layers are. Label noise can be grouped, and is mostly treated based on what is known about the noise generating distribution [?]. Some datasets contain label noise from a known and quantifiable generative distribution, while in other cases, too little or nothing is known about the noise generating distribution to form a hypothesis. Below we discuss the sources and types of label noise in more detail

Sources of Label Noise Label noise is inherent in the data collection and processing life-cycle. Most real world datasets are subjected to a number of label noise sources based on how the data is collected, curated and stored. Label noise in practice broadly stems from: (1) incorrect crowd sourced labels where the annotators are non-experts such as is the case with [Clickworker](#), and [Amazon Mechanical Turk](#). (2) Incorrect expert annotations due to the complexity of the data, as is common in medical fields [?]. (3) Labelling errors introduced by software labelling systems such as [?], scale,, (4) Noised introduced by multiple experts or non experts labelling the same sample differently [?].

Types of Label Noise

3 Deep Active Learning Algorithms For Noisy Labels

In this section we focus on the main contribution of this manuscript: exploring literature on deep active learning algorithms used for vision tasks in the presence of label noise.

3.1 Noise Prior Induced Algorithms

3.2 Noise Robust Deep Active Learning Algorithms

4 Evaluation Datasets and Metrics

4.1 Datasets

4.2 Evaluation Metrics

5 Future Research Directions

6 Conclusion

References

1. Algan, G., Ulusoy, I.: Image classification with deep learning in the presence of noisy labels: A survey. ArXiv **abs/1912.05170** (2021)
2. Arik, S., Pfister, T.: Tabnet: Attentive interpretable tabular learning. In: Proceedings of the AAAI Conference on Artificial Intelligence. vol. 35, pp. 6679–6687 (2021). <https://doi.org/10.1609/aaai.v35i8.16826>
3. Baohua, S., Lin, Y., Wenhan, Z., Michael, L., Patrick, D., Charles, Y., Jason, D.: Supertml: Two-dimensional word embedding for the precognition on structured tabular data. In: 2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops. pp. 2973–2981 (06 2019). <https://doi.org/10.1109/CVPRW.2019.00360>
4. Barnett, S.: Convergence problems with generative adversarial networks (gans). ArXiv **abs/1806.11382** (2018)
5. Breiman, L.: Random forests. Machine Learning Journal **45**, 5–32 (2001). <https://doi.org/10.1023/A:1010933404324>
6. Chen, T., Guestrin, C.: Xgboost: A scalable tree boosting system. In: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. pp. 785–794 (08 2016). <https://doi.org/10.1145/2939672.2939785>
7. Chen, T., Kornblith, S., Norouzi, M., Hinton, G.: A simple framework for contrastive learning of visual representations. In: Proceedings of the 37th International Conference on Machine Learning. pp. 1597–1607. ICML’20, JMLR.org (2020)
8. Chicheng, Z., Kamalika, C.: Active learning from weak and strong labelers. In: Cortes, C., Lawrence, N., Lee, D., Sugiyama, M., Garnett, R. (eds.) Advances in Neural Information Processing Systems. vol. 28. Curran Associates, Inc. (2015), <https://proceedings.neurips.cc/paper/2015/file/eba0dc302bcd9a273f8bbb72be3a687b-Paper.pdf>
9. Cordeiro, F., Carneiro, G.: In: The 33rd SIBGRAPI Conference on Graphics, Patterns and Images. pp. 9–16 (11 2020)
10. C.Shui, F.Zhou, C.Gagn’e, B.Wang: Deep active learning: Unified and principled method for query and training. In: International Conference on Artificial Intelligence and Statistics (2020)
11. Denil, M., Bazzani, L., Larochelle, H., de Freitas, N.: Learning where to attend with deep architectures for image tracking. In: Institute of Electrical and Electronics Engineers, Neural Computation. vol. 24, pp. 2151–2184 (2012)

12. Du, P., Zhao, S., Chen, H., Chai, S., Chen, H., Li, C.: Contrastive coding for active learning under class distribution mismatch. In: IEEE/CVF International Conference on Computer Vision (ICCV). pp. 8907–8916 (2021). <https://doi.org/10.1109/ICCV48922.2021.00880>
13. Geifman, Y., El-Yaniv, R.: Deep active learning over the long tail. ArXiv **abs/1711.00941** (2017)
14. Gonog, L., Zhou, Y.: A review: Generative adversarial networks. In: The 14th IEEE Conference on Industrial Electronics and Applications (ICIEA). pp. 505–510 (2019). <https://doi.org/10.1109/ICIEA.2019.8833686>
15. Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., Bengio, Y.: Generative adversarial networks. In: Ghahramani, Z., Welling, M., Cortes, C., Lawrence, N., Weinberger, K. (eds.) Advances in Neural Information Processing Systems. vol. 27. Curran Associates, Inc. (2014), <https://proceedings.neurips.cc/paper/2014/file/5ca3e9b122f61f8f06494c97b1afccf3-Paper.pdf>
16. Grave, E., Bojanowski, P., Gupta, P., Joulin, A., Mikolov, T.: Learning word vectors for 157 languages. In: Proceedings of the International Conference on Language Resources and Evaluation (LREC 2018) (2018)
17. Har-Peled, S., Roth, D., Zimak, D.: Maximum margin coresets for active and noise tolerant learning. In: International Joint Conferences on Artificial Intelligence (2007)
18. Hochreiter, S., Schmidhuber, J.: Long short-term memory. In: Neural Computation. vol. 9, pp. 1735–1780 (1997)
19. Kolesnikov, A., Dosovitskiy, A., Weissenborn, D., Heigold, G., Uszkoreit, J., Beyer, L., Minderer, M., Dehghani, M., Houlsby, N., Gelly, S., Unterthiner, T., Zhai, X.: An image is worth 16x16 words: Transformers for image recognition at scale. In: 9th International Conference on Learning Representations, ICLR 2021 (2021), <https://openreview.net/forum?id=YicbFdNTTy>
20. Lewis, D., Gale, W.:
21. McCallum, A., Nigam, K.: Employing em and pool-based active learning for text classification. In: International Conference of Machine Learning (1998)
22. Mescheder, L., Geiger, A., Nowozin, S.: Which training methods for gans do actually converge? In: The 35th International Conference on Machine Learning (2018)
23. Novák, A., Laki, L., Novák, B.: CBOW-tag: a modified CBOW algorithm for generating embedding models from annotated corpora. In: Proceedings of the Twelfth Language Resources and Evaluation Conference. pp. 4798–4801. European Language Resources Association (09 2020), <https://aclanthology.org/2020.lrec-1.590>
24. Pennington, J., Socher, R., Manning, C.: Glove: Global vectors for word representation. In: Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP). pp. 1532–1543 (10 2014). <https://doi.org/10.3115/v1/D14-1162>, <http://www.aclweb.org/anthology/D14-1162>
25. Phillips, J.: Coresets and sketches. CoRR **abs/1601.00617** (2016), <http://arxiv.org/abs/1601.00617>
26. Phillips, J., Tai, W.: Near-optimal coresets of kernel density estimates. In: 34th International Symposium on Computational Geometry, SoCG 2018, June 11-14, 2018, Budapest, Hungary. LIPIcs, vol. 99, pp. 66:1–66:13. Schloss Dagstuhl - Leibniz-Zentrum für Informatik (2018). <https://doi.org/10.4230/LIPIcs.SocG.2018.66>, <https://doi.org/10.4230/LIPIcs.SocG.2018.66>

27. Popov, S., Morozov, S., Babenko, A.: Neural oblivious decision ensembles for deep learning on tabular data. ArXiv **abs/1909.06312** (2020)
28. Prokhorenkova, L., Gleb, G., Vorobev, A., Dorogush, A., Gulin, A.: Catboost: unbiased boosting with categorical features. In: Bengio, S., Wallach, H., Larochelle, H., Grauman, K., Cesa-Bianchi, N., Garnett, R. (eds.) Proceedings of the 32nd International Conference on Neural Information Processing Systems. vol. 31, p. 6639–6649. Curran Associates, Inc. (2018)
29. Roman, L., Valeriia, C., Avi, S., Arpit, B., Bruss, C., Tom, G., Andrew, W., Micah, G.: Transfer learning with deep tabular models (06 2022). <https://doi.org/10.48550/arXiv.2206.15306>
30. Rosenblatt, F.: The perceptron: A probabilistic model for information storage and organization in the brain. In: Psychological Review. vol. 65, pp. 386–408 (1958)
31. Schäfl, B., Gruber, L., Bitto-Nemling, A., Hochreiter, S.: Hopular: Modern hopfield networks for tabular data. ArXiv **abs/2206.00664** (2022)
32. See, A., Liu, P., Manning, C.: Get to the point: Summarization with pointer-generator networks. In: Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics. vol. 1, pp. 1073–1083 (2017)
33. Sener, O., Savarese, S.: Active learning for convolutional neural networks: A core-set approach. International Conference on Learning Representations (Poster) (2018)
34. Sinha, S., Ebrahimi, S., Darrell, T.: Ieee/cvf international conference on computer vision (iccv) (03 2019)
35. Srinivas1, A., Lin, T., Parmar, N., Shlens, J., Abbeel, P., Vaswani, A.: Bottleneck transformers for visual recognition. In: 2021 Conference on Computer Vision and Pattern Recognition (2021)
36. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A., Kaiser, L., Polosukhin, I.: Attention is all you need. In: Advances in Neural Information Processing Systems. vol. 30 (2017), <https://arxiv.org/pdf/1706.03762.pdf>
37. Wang, C., Singla, A., Chen, Y.: Teaching an active learner with contrastive examples. In: Advances in Neural Information Processing Systems (2021)
38. Wei, J., Zhu, Z., Cheng, H., Liu, T., Niu, G., Liu, Y.: Learning with noisy labels revisited: A study using real-world human annotations. 10th International Conference on Learning Representations (2022)